

nir

June 2, 2022

0.1

1. .
2. . , .
3. , . . . ,
4. . . .
5. . .
6. . ,
7. .
8. (baseline) .
9. - .
10. 8 . baseline- .
11. .

0.2 1.

Diabetes Health Indicators Dataset
(<https://www.kaggle.com/datasets/alexteboul/diabetes-health-indicators-dataset>).

csv :

- Diabetes_binary - 1 - / 0 - .
- HighBP - 1 - / 0 - .
- HighChol - 1 - / 0 -
- CholCheck - 1 - 5- / 0 -
- BMI -
- Smoker - 100 ? - 1/ - 0
- Stroke - ? - 1/ - 0.
- HeartDiseaseorAttack - () () 0 = 1 =
- PhysActivity - 30 , 0 = 1 =
- Fruits - 1 0 = 1 =
- Veggies - 1 0 = 1 =
- HvyAlcoholConsumption - (, 14 ,
7) 0 =
- AnyHealthcare - , , HMO
. . 0 = 1 =

- NoDocbcCost - 12 , , -
? 0 = 1 =
- GenHlth - , : 1 5 1 = 2 = 3 = 4
= 5 =
- MentHlth - , , , ,
- PhysHlth - , , , ,
30
- DiffWalk - ? 0 = 1 =
- Sex - 0 = 1 =
- Age - 13- (_AGEG5YR, .) 1 = 18-24 9 = 60-64 13 =
80
- Education - (EDUCA .)
- Income - (INCOME2 .)

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from collections import Counter

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
↳ GradientBoostingClassifier
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score,
↳ ConfusionMatrixDisplay, precision_score, recall_score, f1_score,
↳ classification_report, roc_curve, plot_roc_curve, auc,
↳ precision_recall_curve, plot_precision_recall_curve, average_precision_score
from sklearn.model_selection import GridSearchCV

import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: data = pd.read_csv('./diabetes_binary_5050split_health_indicators_BRFSS2015.
↳ csv')
```

```
[ ]: data.head()
```

```
[ ]:      Diabetes_binary  HighBP  HighChol  CholCheck  BMI  Smoker  Stroke  \
0                0.0    1.0    0.0    1.0  26.0    0.0    0.0
1                0.0    1.0    1.0    1.0  26.0    1.0    1.0
2                0.0    0.0    0.0    1.0  26.0    0.0    0.0
3                0.0    1.0    1.0    1.0  28.0    1.0    0.0
4                0.0    0.0    0.0    1.0  29.0    1.0    0.0

      HeartDiseaseorAttack  PhysActivity  Fruits  ...  AnyHealthcare  \
0                0.0            1.0    0.0  ...            1.0
1                0.0            0.0    1.0  ...            1.0
2                0.0            1.0    1.0  ...            1.0
3                0.0            1.0    1.0  ...            1.0
4                0.0            1.0    1.0  ...            1.0

      NoDocbcCost  GenHlth  MentHlth  PhysHlth  DiffWalk  Sex  Age  Education  \
0                0.0    3.0    5.0    30.0    0.0  1.0  4.0    6.0
1                0.0    3.0    0.0    0.0    0.0  1.0  12.0   6.0
2                0.0    1.0    0.0   10.0    0.0  1.0  13.0   6.0
3                0.0    3.0    0.0    3.0    0.0  1.0  11.0   6.0
4                0.0    2.0    0.0    0.0    0.0  0.0  8.0   5.0

      Income
0      8.0
1      8.0
2      8.0
3      8.0
4      8.0

[5 rows x 22 columns]
```

```
[ ]: print(f'          {data.shape[0]}          {data.shape[1]}          .')
```

```
70692      22      .
```

```
[ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 70692 entries, 0 to 70691
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Diabetes_binary       70692 non-null  float64
1   HighBP                70692 non-null  float64
2   HighChol              70692 non-null  float64
3   CholCheck             70692 non-null  float64
4   BMI                   70692 non-null  float64
```

```

5   Smoker                70692 non-null float64
6   Stroke                70692 non-null float64
7   HeartDiseaseorAttack  70692 non-null float64
8   PhysActivity          70692 non-null float64
9   Fruits                70692 non-null float64
10  Veggies               70692 non-null float64
11  HvyAlcoholConsump     70692 non-null float64
12  AnyHealthcare         70692 non-null float64
13  NoDocbcCost           70692 non-null float64
14  GenHlth               70692 non-null float64
15  MentHlth              70692 non-null float64
16  PhysHlth              70692 non-null float64
17  DiffWalk              70692 non-null float64
18  Sex                   70692 non-null float64
19  Age                   70692 non-null float64
20  Education             70692 non-null float64
21  Income                70692 non-null float64
dtypes: float64(22)
memory usage: 11.9 MB

```

```
[ ]: data = data.astype('int')
```

0.3 2.

```
[ ]: data.isnull().sum()
```

```

[ ]: Diabetes_binary      0
    HighBP                0
    HighChol              0
    CholCheck             0
    BMI                   0
    Smoker                 0
    Stroke                 0
    HeartDiseaseorAttack  0
    PhysActivity           0
    Fruits                 0
    Veggies                0
    HvyAlcoholConsump     0
    AnyHealthcare         0
    NoDocbcCost           0
    GenHlth                0
    MentHlth              0
    PhysHlth              0
    DiffWalk              0
    Sex                    0
    Age                    0
    Education              0

```


```
Income
dtype: int64
```

```
[ ]: total = data.shape[0]
class_0, class_1 = data['Diabetes_binary'].value_counts()
print('    0    {},    1    {}'.format(round(class_0 / total, 4)*100, round(class_1 / total, 4)*100))
```

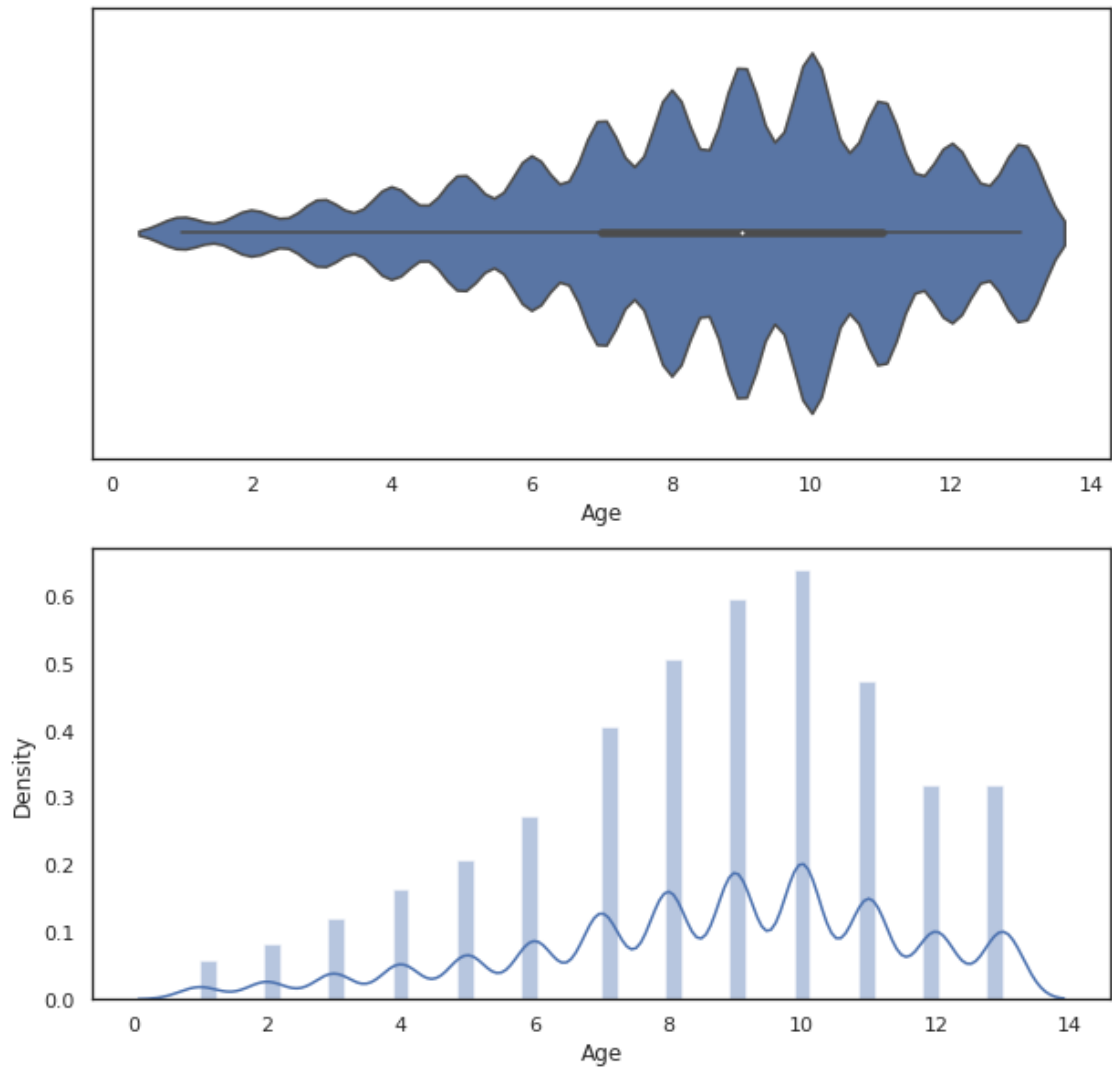
```
0    50.0%,    1    50.0%.
```


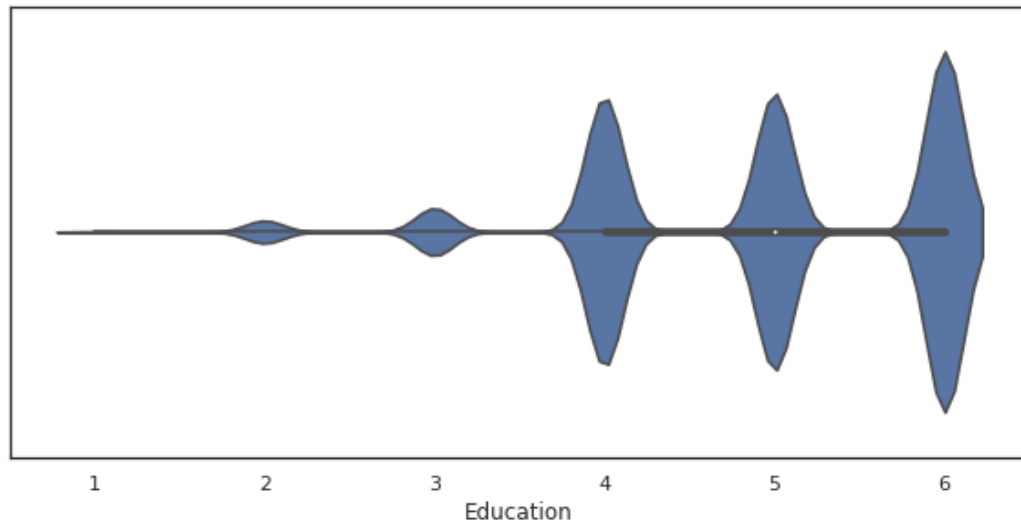
```
[ ]: mini_data = data.sample(frac=1)
mini_data = mini_data[:500]
sns.pairplot(mini_data, hue='Diabetes_binary')
```



```
[ ]: <seaborn.axisgrid.PairGrid at 0x7f32494584c0>
```





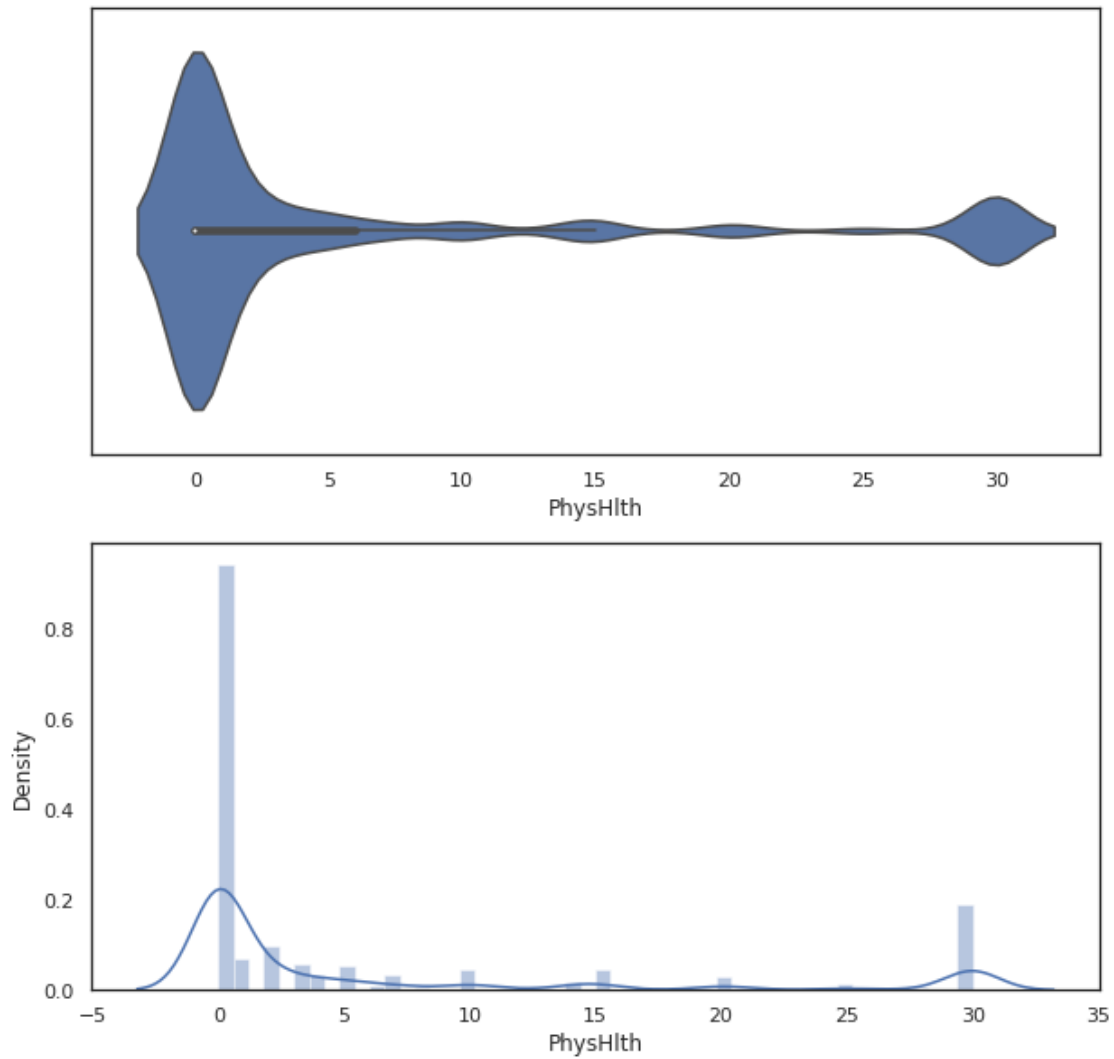
```
[ ]: for col in ['Age', 'Education', 'Income', 'BMI', 'PhysHlth']:  
    fig, ax = plt.subplots(2, 1, figsize=(10,10))  
    sns.violinplot(ax=ax[0], x=data[col])  
    sns.distplot(data[col], ax=ax[1])
```












0.3.1 3. .

```
[ ]: fig, ax = plt.subplots(1, 1, sharex='col', sharey='row', figsize=(20,10))
     sns.heatmap(data.corr(), annot=True, fmt='.2f')
```

```
[ ]: <AxesSubplot:>
```



```
[ ]: def correlation_fun(ds,threshold):
    corr_col=set()
    corr_mat=ds.corr()
    for i in range(len(corr_mat.columns)):
        for j in range(i):
            if abs(corr_mat.iloc[i,j])>threshold:
                colname=corr_mat.columns[i]
                corr_col.add(colname)
    return corr_col
```

```
[ ]: threshold = 0.5

correlation_fun(data.drop("Diabetes_binary",axis=1),threshold)
```

```
[ ]: {'PhysHlth'}
```

Sex, AnyHealthcare, NoDocbcCost, Fruits, PhysHlth.

```
[ ]: drop_columns = ['Sex','AnyHealthcare', 'NoDocbcCost', 'Fruits', 'PhysHlth']
```

```
[ ]: data=data.drop(drop_columns,axis=1)
```

```
[ ]: data.head()
```

```
[ ]: Diabetes_binary HighBP HighChol CholCheck BMI Smoker Stroke \
0      0      1      0      1  26      0      0
1      0      1      1      1  26      1      1
2      0      0      0      1  26      0      0
3      0      1      1      1  28      1      0
4      0      0      0      1  29      1      0

HeartDiseaseorAttack PhysActivity Veggies HvyAlcoholConsump GenHlth \
0      0      1      1      0      3
1      0      0      0      0      3
2      0      1      1      0      1
3      0      1      1      0      3
4      0      1      1      0      2

MentHlth DiffWalk Age Education Income
0      5      0   4      6      8
1      0      0  12      6      8
2      0      0  13      6      8
3      0      0  11      6      8
4      0      0   8      5      8
```

```
[ ]: data.shape
```

```
[ ]: (70692, 17)
```

0.3.2

,
 precision:
 , “accuracy”.

$$precision = \frac{TP}{TP + FP}$$

precision_score.
 recall ():

$$recall = \frac{TP}{TP + FN}$$

recall_score.
 F1-

, precision recall F_β ,
precision recall:

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

β .
F1- (F-) $\beta=1$:

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

f1_score.

ROC AUC

:

$$\text{truePR} = \frac{TP}{TP + FN}$$

True Positive Rate, . recall.

$$\text{falsePR} = \frac{FP}{FP + TN}$$

False Positive Rate, . .

ROC- (0,0)-(0,1)-(1,1), .

,

- ROC AUC (Area Under the Receiver Operating Characteristic Curve).

ROC AUC roc_auc_score.

0.3.3

.

: - - - -

0.3.4

.

```
[ ]: X = data.drop('Diabetes_binary', axis=1)
      Y = data['Diabetes_binary']

[ ]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
      ↪random_state=1)

[ ]: X_train.shape
```

```
[ ]: (56553, 16)
```

```
[ ]: X_test.shape
```

```
[ ]: (14139, 16)
```

```
[ ]: sc = StandardScaler()

X_train = sc.fit_transform(X_train)

X_test = sc.transform(X_test)
```

0.3.5

0.3.6

```
[ ]: models = { 'LogisticRegression': LogisticRegression(),
                'KNearestNeighbors': KNeighborsClassifier(n_neighbors=5),
                'DecisionTree': DecisionTreeClassifier(),
                'RandomForest': RandomForestClassifier(),
                'GradientBoost': GradientBoostingClassifier()}

accuracies = {}
```

```
[ ]: def DrawGraphics(Y_test, y_pred):
    print("*****")
    print(model_name)
    print("*****")
    print(classification_report(Y_test, y_pred))
    print(f'ROC AUC score: {roc_auc_score(Y_test, y_prob)}')
    print('Accuracy Score: ', accuracy_score(Y_test, y_pred))

    plt.figure(figsize = (6, 6))
    sns.heatmap(cm, cmap = 'Blues', annot = True, fmt = 'd', linewidths = 5,
    ↪ cbar = False, annot_kws = {'fontsize': 15},
    yticklabels = ['Healthy', 'Diabetic'], xticklabels = ['Predicted_
    ↪ Healthy', 'Predicted Diabetic'])
    plt.yticks(rotation = 0)
    plt.show()

    false_positive_rate, true_positive_rate, thresholds = roc_curve(Y_test,
    ↪ y_prob)
    roc_auc = auc(false_positive_rate, true_positive_rate)

    sns.set_theme(style = 'white')
    plt.figure(figsize = (6, 6))
```

```

plt.plot(false_positive_rate,true_positive_rate, color = '#b01717', label = 'AUC = %0.3f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], linestyle = '--', color = '#174ab0')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.title('ROC AUC Curve')
plt.legend()
plt.show()

```

```

[ ]: for model_name, model in models.items():
    model.fit(X_train, Y_train)
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:,-1]
    cm = confusion_matrix(Y_test, y_pred)

    DrawGraphics(Y_test, y_pred)

    acc = accuracy_score(Y_test, y_pred)*100
    accuracies[model_name] = acc

```


LogisticRegression

	precision	recall	f1-score	support
0	0.76	0.73	0.75	7141
1	0.74	0.77	0.75	6998
accuracy			0.75	14139
macro avg	0.75	0.75	0.75	14139
weighted avg	0.75	0.75	0.75	14139

ROC AUC score: 0.8276305683433108

Accuracy Score: 0.7484970648560718

Healthy	5229	1912
Diabetic	1644	5354
	Predicted Healthy	Predicted Diabetic




KNearestNeighbors

	precision	recall	f1-score	support
0	0.73	0.69	0.71	7141
1	0.70	0.74	0.72	6998
accuracy			0.71	14139
macro avg	0.72	0.71	0.71	14139
weighted avg	0.72	0.71	0.71	14139

ROC AUC score: 0.7720235069063084

Accuracy Score: 0.7146191385529387

Healthy	4922	2219
Diabetic	1816	5182
	Predicted Healthy	Predicted Diabetic




DecisionTree

	precision	recall	f1-score	support
0	0.66	0.68	0.67	7141
1	0.66	0.64	0.65	6998
accuracy			0.66	14139
macro avg	0.66	0.66	0.66	14139
weighted avg	0.66	0.66	0.66	14139

ROC AUC score: 0.6607924047677376

Accuracy Score: 0.6619987269255252

Healthy	4873	2268
Diabetic	2511	4487
	Predicted Healthy	Predicted Diabetic




RandomForest

	precision	recall	f1-score	support
0	0.75	0.70	0.72	7141
1	0.71	0.76	0.74	6998
accuracy			0.73	14139
macro avg	0.73	0.73	0.73	14139
weighted avg	0.73	0.73	0.73	14139

ROC AUC score: 0.7999799910823342

Accuracy Score: 0.7293302213735059

Healthy	4964	2177
Diabetic	1650	5348
	Predicted Healthy	Predicted Diabetic




GradientBoost

	precision	recall	f1-score	support
0	0.78	0.72	0.75	7141
1	0.73	0.79	0.76	6998
accuracy			0.75	14139
macro avg	0.76	0.75	0.75	14139
weighted avg	0.76	0.75	0.75	14139

ROC AUC score: 0.8347393171610158

Accuracy Score: 0.7538015418346418

Healthy	5111	2030
Diabetic	1451	5547
	Predicted Healthy	Predicted Diabetic



```
[ ]: n_range_list = list(range(0,250,50))
      n_range_list[0] = 1
      n_range_list
```

```
[ ]: [1, 50, 100, 150, 200]
```

```
[ ]: grid_models = [(LogisticRegression(),[{ 'C': [0.25,0.5,0.75,1], 'random_state':
      ↳ [0]}])),
                    (KNeighborsClassifier(),[{ 'n_neighbors': n_range_list}]),
                    (DecisionTreeClassifier(),[{ 'criterion':
      ↳ ['gini','entropy'], 'random_state': [0]}])),
                    (RandomForestClassifier(),[{ 'n_estimators':
      ↳ n_range_list, 'criterion': ['gini','entropy'], 'random_state': [0]}])),
```

```

        (GradientBoostingClassifier(),[{'n_estimators':
↪n_range_list, 'criterion':['friedman_mse','mse'],'loss':
↪['deviance','exponential'],'learning_rate':[0.1, 0.5, 0.8, 1],'random_state':
↪[0]}]))

```

```

[ ]: for i,j in grid_models:
    grid = GridSearchCV(estimator=i,param_grid = j, scoring = 'accuracy',cv=2)
    grid.fit(X_train, Y_train)
    best_accuracy = grid.best_score_
    best_param = grid.best_params_
    print('{:}\nBest Accuracy : {:.2f}%'.format(i,best_accuracy*100))
    print('Best Parameters : ',best_param)
    print('')
    print('-----')
    print('')

```

```

LogisticRegression():
Best Accuracy : 74.67%
Best Parameters : {'C': 0.5, 'random_state': 0}

```

```

KNeighborsClassifier():
Best Accuracy : 74.08%
Best Parameters : {'n_neighbors': 150}

```

```

DecisionTreeClassifier():
Best Accuracy : 65.79%
Best Parameters : {'criterion': 'entropy', 'random_state': 0}

```

```

RandomForestClassifier():
Best Accuracy : 72.84%
Best Parameters : {'criterion': 'entropy', 'n_estimators': 200, 'random_state':
0}

```

```

GradientBoostingClassifier():
Best Accuracy : 75.01%
Best Parameters : {'criterion': 'friedman_mse', 'learning_rate': 0.1, 'loss':
'exponential', 'n_estimators': 200, 'random_state': 0}

```

```
[ ]: params_models = { 'LogisticRegression': LogisticRegression(C = 0.5,
↳random_state= 0),
                        'KNearestNeighbors': KNeighborsClassifier(n_neighbors=150),
                        'DecisionTree':
↳DecisionTreeClassifier(criterion='entropy',random_state=0),
                        'RandomForest':
↳RandomForestClassifier(criterion='gini',n_estimators=200,random_state=0),
                        'GradientBoost':
↳GradientBoostingClassifier(criterion='friedman_mse',learning_rate=0.
↳1,loss='exponential',n_estimators=200,random_state=0)}

params_accuracies = {}
params_precision = {}
params_recall = {}
params_f1 = {}
params_roc_auc = {}
```

```
[ ]: for model_name, model in params_models.items():
    model.fit(X_train, Y_train)
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[: ,1]
    cm = confusion_matrix(Y_test, y_pred)

    DrawGraphics(Y_test, y_pred)

    param_acc = accuracy_score(Y_test, y_pred)*100
    params_accuracies[model_name] = param_acc
    params_precision[model_name] = precision_score(Y_test, y_pred)
    params_recall[model_name] = recall_score(Y_test, y_pred)
    params_f1[model_name] = f1_score(Y_test, y_pred)
    params_roc_auc[model_name] = roc_auc_score(Y_test,y_pred)
```


LogisticRegression

	precision	recall	f1-score	support
0	0.76	0.73	0.75	7141
1	0.74	0.77	0.75	6998
accuracy			0.75	14139
macro avg	0.75	0.75	0.75	14139
weighted avg	0.75	0.75	0.75	14139

ROC AUC score: 0.827631468834655

Accuracy Score: 0.7484970648560718

Healthy	5229	1912
Diabetic	1644	5354
	Predicted Healthy	Predicted Diabetic




KNearestNeighbors

	precision	recall	f1-score	support
0	0.77	0.69	0.73	7141
1	0.72	0.79	0.75	6998
accuracy			0.74	14139
macro avg	0.75	0.74	0.74	14139
weighted avg	0.75	0.74	0.74	14139

ROC AUC score: 0.8216396794747087

Accuracy Score: 0.7423438715609307

Healthy	4942	2199
Diabetic	1444	5554
	Predicted Healthy	Predicted Diabetic




DecisionTree

	precision	recall	f1-score	support
0	0.66	0.69	0.68	7141
1	0.67	0.64	0.66	6998
accuracy			0.67	14139
macro avg	0.67	0.67	0.67	14139
weighted avg	0.67	0.67	0.67	14139

ROC AUC score: 0.6652005200117392

Accuracy Score: 0.6661715821486668

Healthy	4912	2229
Diabetic	2491	4507
	Predicted Healthy	Predicted Diabetic




RandomForest

	precision	recall	f1-score	support
0	0.76	0.69	0.72	7141
1	0.71	0.77	0.74	6998
accuracy			0.73	14139
macro avg	0.73	0.73	0.73	14139
weighted avg	0.73	0.73	0.73	14139

ROC AUC score: 0.800058894134996

Accuracy Score: 0.7318763703232195

Healthy	4956	2185
Diabetic	1606	5392
	Predicted Healthy	Predicted Diabetic




GradientBoost

	precision	recall	f1-score	support
0	0.78	0.71	0.75	7141
1	0.73	0.80	0.76	6998
accuracy			0.76	14139
macro avg	0.76	0.76	0.76	14139
weighted avg	0.76	0.76	0.76	14139

ROC AUC score: 0.8349524634621635


Accuracy Score: 0.7564891435037838

Healthy	5094	2047
Diabetic	1396	5602
	Predicted Healthy	Predicted Diabetic




```
[ ]: colors = ["purple", "green", "orange",
               ↪ "magenta", "#CFC60E", "#0FBBAE", '#417D7A', '#066163', '#4D4C7D']

sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,100,10))
plt.ylabel("Accuracy %")
plt.xlabel("Algorithms")
sns.barplot(x=list(accuracies.keys()), y=list(accuracies.values()),
            ↪ palette=colors)
plt.show()
```



```
[ ]: colors = ["purple", "green", "orange", "\u2192"
               "magenta", "#CFC60E", "#0FBBAE", '#417D7A', '#066163', '#4D4C7D']


sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,100,10))
plt.ylabel("Accuracy %")
plt.xlabel("Algorithms")
sns.barplot(x=list(params_accuracies.keys()), y=list(params_accuracies.
               \u2192values()), palette=colors)
plt.show()
```



```
[ ]: colors = ["purple", "green", "orange", "\u2192"
               "magenta", "#CFC60E", "#0FBBAE", '#417D7A', '#066163', '#4D4C7D']

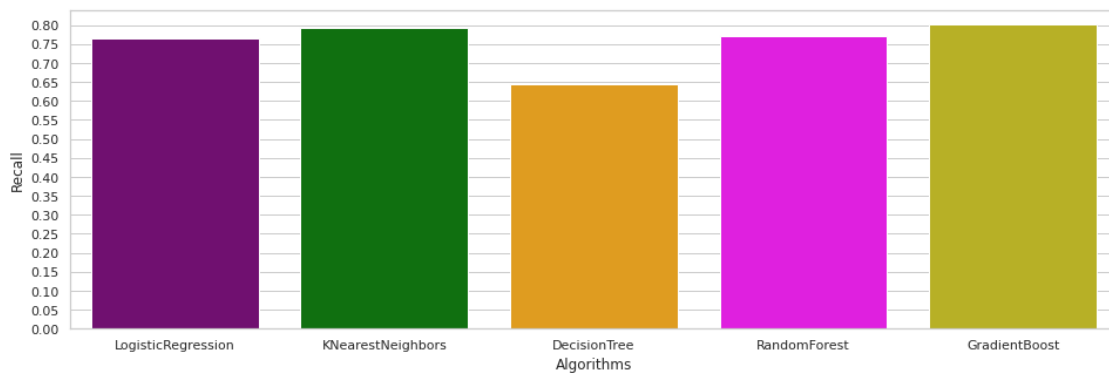
sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,1,0.05))
plt.ylabel("Precision")
```

```
plt.xlabel("Algorithms")
sns.barplot(x=list(params_accuracies.keys()), y=list(params_precision.
↪ values()), palette=colors)
plt.show()
```



```
[ ]: colors = ["purple", "green", "orange", ↪
↪ "magenta", "#CFC60E", "#0FBBAE", '#417D7A', '#066163', '#4D4C7D']

sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,1,0.05))
plt.ylabel("Recall")
plt.xlabel("Algorithms")
sns.barplot(x=list(params_accuracies.keys()), y=list(params_recall.values()), ↪
↪ palette=colors)
plt.show()
```




```
[ ]: colors = ["purple", "green", "orange", ↪
↪ "magenta", "#CFC60E", "#0FBBAE", '#417D7A', '#066163', '#4D4C7D']
```

```

sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,1,0.05))
plt.ylabel("f1")
plt.xlabel("Algorithms")
sns.barplot(x=list(params_accuracies.keys()), y=list(params_f1.values()),
            palette=colors)
plt.show()

```




```

[ ]: colors = ["purple", "green", "orange",
               "magenta", "#CFC60E", "#0FBBAE", '#417D7A', '#066163', '#4D4C7D']

sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,1,0.05))
plt.ylabel("ROC_AUC")
plt.xlabel("Algorithms")
sns.barplot(x=list(params_accuracies.keys()), y=list(params_roc_auc.values()),
            palette=colors)
plt.show()

```



:

.